

Practical 10

Task1: Investigate inode structures using commands such as `ls -li`, `stat`, and `find`. Create hard links and symbolic links and explain their impact on inode allocation.

Step 1 — Create the original file

```
bhavan@DESKTOP-BLT5L3R:~$ cd ~/OSSP
mkdir Week10
cd Week10
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ pwd
/home/bhavan/OSSP/Week10
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ ls
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ echo "This is the original file for
Week 10 Task 1." > original.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ cat original.txt
This is the original file for Week 10 Task 1.
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ ls -li original.txt
-rw-r--r-- 1 bhavan bhavan 46 Sep 22 05:39 original.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$
```

Step 2 — Check the inode number

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ ls -li original.txt
1785 original.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$
```

Your original file has:

- **Inode number:** 1785
- **Link count:** 1
- **File size:** 46 bytes

Step 3 — Use `stat` to investigate the inode

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ stat original.txt
  File: original.txt
  Size: 46          Blocks: 8          IO Block: 4096   regular file
Device: 8,48      Inode: 1785         Links: 1
Access: (0644/-rw-r--r--)  Uid: ( 1000/bhavan.)   Gid: ( 1000/bhavan.)
Access: 2026-09-22 05:39:36.168572457 +0000
Modify: 2026-09-22 05:39:31.096862456 +0000
Change: 2026-09-22 05:39:31.096862456 +0000
 Birth: 2026-09-22 05:39:31.096862456 +0000
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$
```

Step 4 — Find the file using its inode

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ find . -inum 1785
./original.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$
```

Step 5 — Create a hard link

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ ln original.txt hardlink.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ ls -li original.txt hardlink.txt
1785 -rw-r--r-- 2 bhavan bhavan 46 Sep 22 05:39 hardlink.txt
1785 -rw-r--r-- 2 bhavan bhavan 46 Sep 22 05:39 original.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ find . -inum 1785
./hardlink.txt
./original.txt
```

Step 6 — Create a symbolic link

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ ln -s original.txt symlink.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ ls -li original.txt hardlink.txt sym
link.txt
1785 -rw-r--r-- 2 bhavan bhavan 46 Sep 22 05:39 hardlink.txt
1785 -rw-r--r-- 2 bhavan bhavan 46 Sep 22 05:39 original.txt
12935 lrwxrwxrwx 1 bhavan bhavan 12 Sep 22 05:46 symlink.txt -> original
.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ |
```

Step 7 — Check the symbolic link using stat

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ stat symlink.txt
  File: symlink.txt -> original.txt
  Size: 12          Blocks: 0          IO Block: 4096   symbolic link
Device: 8,48      Inode: 12935         Links: 1
Access: (0777/lrwxrwxrwx)  Uid: ( 1000/bhavan.)   Gid: ( 1000/bhavan.)
Access: 2026-09-22 05:46:59.999498628 +0000
Modify: 2026-09-22 05:46:53.423509843 +0000
Change: 2026-09-22 05:46:53.423509843 +0000
 Birth: 2026-09-22 05:46:53.423509843 +0000
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ stat -L symlink.txt
  File: symlink.txt
  Size: 46          Blocks: 8          IO Block: 4096   regular file
Device: 8,48      Inode: 1785         Links: 2
Access: (0644/-rw-r--r--)  Uid: ( 1000/bhavan.)   Gid: ( 1000/bhavan.)
Access: 2026-09-22 05:39:36.168572457 +0000
Modify: 2026-09-22 05:39:31.096862456 +0000
Change: 2026-09-22 05:45:34.620200224 +0000
 Birth: 2026-09-22 05:39:31.096862456 +0000
```

Step 8 — Find the symbolic link's inode

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ ls -li symlink.txt
12935 symlink.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ find . -inum 1786
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ find . -inum 12935
./symlink.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ stat symlink.txt
  File: symlink.txt -> original.txt
  Size: 12          Blocks: 0          IO Block: 4096   symbolic link
Device: 8,48      Inode: 12935         Links: 1
Access: (0777/lrwxrwxrwx)  Uid: ( 1000/bhavan.)   Gid: ( 1000/bhavan.)
Access: 2026-09-22 05:46:59.999498628 +0000
Modify: 2026-09-22 05:46:53.423509843 +0000
Change: 2026-09-22 05:46:53.423509843 +0000
 Birth: 2026-09-22 05:46:53.423509843 +0000
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ |
```

Feature	Hard Link	Symbolic Link
Inode	Same inode as original	Separate inode
File data	Shared	Target's data is accessed
Link count	Increases original inode's link count	Does not increase target's link count
Can cross file systems?	No	Yes
If original is deleted	Data remains accessible	Link becomes broken
Creation	ln file link	ln -s file link

Observation:

Hard links **do not allocate a new inode** for the linked file; they create another directory entry pointing to the existing inode. Symbolic links **allocate a new inode** because the symbolic link itself is a separate file containing the path to the target. Thus, hard links conserve inode allocation, while symbolic links require an additional inode.

Task2: Design a program using mmap() to perform file reading and writing. Compare memory-mapped I/O with traditional read/write operations in terms of efficiency and Implementation complexity .

Step 1 – Create input file

nano mmap_file.txt

```
GNU nano 7.2      mmap_file.txt *
This is a file used to demonstrate memory mapped I/O.
The mmap() system call allows the file to be accessed through memory.|
```

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10$ cd ~/OSSP
mkdir Week10_Task2
cd Week10_Task2
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ pwd
/home/bhavan/OSSP/Week10_Task2
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ nano mmap_file.txt
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ cat mmap_file.txt
This is a file used to demonstrate memory mapped I/O.
The mmap() system call allows the file to be accessed through memory.
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ |
```

Step 2: Create the mmap() program

nano mmap_file.c

```
GNU nano 7.2                                mmap_file.c *
#include <stdio.h>
#include <fcntl.h>
#include <unistd.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <string.h>

int main() {
    int fd;
    struct stat st;
    char *data;

    /* Open the file for reading and writing */
    fd = open("mmap_file.txt", O_RDWR);

    if (fd == -1) {
        perror("Error opening file");
        return 1;
    }

    /* Get file size */
    if (fstat(fd, &st) == -1) {
        perror("Error getting file information");
        close(fd);
        return 1;
    }

    /* Map the file into memory */
    data = mmap(NULL, st.st_size,
                PROT_READ | PROT_WRITE,
                MAP_SHARED, fd, 0);

    if (data == MAP_FAILED) {
        perror("Error mapping file");
        close(fd);
        return 1;
    }

    /* Read file contents through memory */
    printf("Original file contents:\n");
    printf("%s\n", data);

    /* Write new content through mapped memory */
    strcpy(data, "File modified using mmap() memory-mapped I/O.\n");

    /* Synchronize changes with the file */
    msync(data, st.st_size, MS_SYNC);

    /* Remove memory mapping */
    munmap(data, st.st_size);

    close(fd);

    printf("File modified successfully using mmap().\n");

    return 0;
}
```

Step 3: Compile the mmap() program

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ nano mmap_file.c
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ gcc mmap_file.c -o mmap_file
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ ./mmap_file
Original file contents:
This is a file used to demonstrate memory mapped I/O.
The mmap() system call allows the file to be accessed through memory.

File modified successfully using mmap().
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ cat mmap_file.txt
File modified using mmap() memory-mapped I/O.
d I/O.
The mmap() system call allows the file to be accessed through memory.
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ |
```

Step 4: Create the traditional read() / write() program

nano traditional_io.c

```
GNU nano 7.2                                traditional_io.c *
#include <stdio.h>
#include <fcntl.h>
#include <unistd.h>
#include <string.h>

int main() {
    int fd;
    char buffer[200];

    /* Open the file for reading and writing */
    fd = open("mmap_file.txt", O_RDWR);

    if (fd == -1) {
        perror("Error opening file");
        return 1;
    }

    /* Read file contents into buffer */
    int bytesRead = read(fd, buffer, sizeof(buffer) - 1);

    if (bytesRead == -1) {
        perror("Error reading file");
        close(fd);
        return 1;
    }

    buffer[bytesRead] = '\0';

    printf("Original file contents:\n");
    printf("%s\n", buffer);

    /* Move file position to beginning */
    lseek(fd, 0, SEEK_SET);

    /* Write new content */
    char message[] = "File modified using read() and write().\n";

    write(fd, message, strlen(message));

    close(fd);

    printf("File modified successfully using read/write I/O.\n");

    return 0;
}
```

Step 5: Compile and run the traditional I/O program

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ nano traditional_io.c
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ gcc traditional_io.c -o traditional_io
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ ./traditional_io
Original file contents:
File modified using mmap() memory-mapped I/O.

File modified successfully using read/write I/O.
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ cat mmap_file.txt
File modified using read() and write().
I/O.
d I/O.
The mmap() system call allows the file to be accessed through memory.
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ |
```

Step 6: Compare the execution time

```
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ time ./mmap_file
Original file contents:
File modified using read() and write().
I/O.

File modified successfully using mmap().

real    0m0.008s
user    0m0.001s
sys     0m0.004s
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ time ./traditional_io
Original file contents:
File modified using mmap() memory-mapped I/O.

File modified successfully using read/write I/O.

real    0m0.003s
user    0m0.000s
sys     0m0.002s
bhavan@DESKTOP-BLT5L3R:~/OSSP/Week10_Task2$ |
```

Observation:

The file was successfully read and modified using both memory-mapped I/O and traditional read()/write() operations. The mmap() method accessed the file through a memory mapping, while the traditional method used a buffer for transferring data. Execution time was measured using the time command. For the small test file, the difference in execution time was very small.

Result:

A file reading and writing program using mmap() was successfully implemented. It was compared with traditional read()/write() operations. mmap() provides direct memory-based access and can be useful for large or repeatedly accessed files, while read()/write() is simpler to implement and suitable for straightforward file operations.